

HW1: Classification

Jiafeng Chen Francisco Rivera

February 19, 2019

1 Introduction

In this write-up, our main focus is language modeling. That is, given words in a sentence, can we predict the word that follows? We implemented a trigram model, a embedding neural network model, an LSTM, and a few extensions—including warm-starting from pretrained embeddings, ensemble models, and multi-head attention decoder.

2 Problem Description

3 Model and Algorithms

4 Experiments

5 Conclusion

References