

HW1: Classification

Jiafeng Chen Francisco Rivera

February 19, 2019

1 Introduction

In this write-up, our main focus is language modeling. That is, given words in a sentence, can we predict the word that follows? We implemented a **Trigram model**, an embedding neural network model, an **LSTM**, and a few extensions—including pre-trained embeddings, ensemble models, and multi-head attention decoders.

2 Problem Description

To tackle language modeling, we start with sequences of words $w \in \mathcal{V}$ in some vocabulary \mathcal{V} and aim to predict the last word in the sequence which we cannot observe. We can do this probabilistically by attempting to estimate,

$$p(w_t \mid w_1, \dots, w_{t-1}) \tag{1}$$

that is, the conditional distribution over the last word conditional on the words leading up to it.

In particular, there will be a special word, $w_{\text{unk}} \in \mathcal{V}$ which represent an unknown word; we use this whenever we encounter a token we have not previously seen.

In some models, we represent words with dense embeddings. That is, each word gets assigned a vector $v \in \mathbb{R}^d$ where d is the embedding dimension. These

embeddings are trained as part of the model, but can also be initialized to pre-trained values.

3 Model and Algorithms

3.1 Trigram model

In our trigram model, we aim to estimate the probability written in Equation 1. This conditional probability is intractable itself because it's likely that we've never seen the exact sequence of words w_1, \dots, w_{t-1} . However, we can gain tractability by dropping words toward the beginning of the sequence, hoping that they don't affect the probability too much. That is, we hope that,

$$p(w_t \mid w_1, \dots, w_{t-1}) \stackrel{?}{\approx} p(w_t \mid w_{t-2}, w_{t-1}).$$

Having replaced our first probability with a simpler one which conditions on less information, we can estimate the latter by its empirical sample estimator. In other words, we can take all the times in our training set when we've seen words w_{t-2}, w_{t-1} adjacent to each other, and consider the empirical distribution of the word that follows them. We represent this sample approximation as \hat{p} and write,

$$p(w_t \mid w_{t-2}, w_{t-1}) \approx \hat{p}(w_t \mid w_{t-2}, w_{t-1}).$$

By doing this, we've solved most of the intractability of conditioning on the entire sentence w_1, \dots, w_{t-1} , but we still have some of the same problems. Namely, it's possible that in our training set, we either haven't seen words w_{t-2} and w_{t-1} together before, or we've seen them only a very small number of times such that the empirical probability distribution becomes a poor approximation. (To avoid division by zero errors, we adopt the convention that empirical probabilities are all 0 if we haven't seen the words being conditioned on before.) We can fix this by also considering the probabilities,

$$p(w_t) \text{ and } p(w_t \mid w_{t-1})$$

which give us the unconditional probability of a word and the probability conditional on only the previous word. These have the benefit of being more tractable to estimate and the drawback of losing information. In the end, we calculate a blend of these three approximations:

$$\alpha_1 \hat{p}(w_t \mid w_{t-2}, w_{t-1}) + \alpha_2 \hat{p}(w_t \mid w_{t-1}) + (1 - \alpha_1 - \alpha_2) \hat{p}(w_t).$$

Training the weights (α_1, α_2) loads up most of our weight on α_1 which suggests the latter two probabilities are better used as “tie-breakers” when conditioning on the previous bi-gram yields a small number of possibilities. In our final model, we use $(\alpha_1, \alpha_2) = (0.9, 0.05)$.

3.2 Neural Network Language Model

Following [Bengio et al. \(2003\)](#), we implement a neural network language model (NNLM). We model (1) by first assuming limited dependence:

$$p(w_i \mid w_1, \dots, w_{i-1}) = p(w_i \mid w_{i-1}, \dots, w_{i-k}),$$

i.e., the current word only depends on the past k words, a useful restriction motivated by n-gram models. Next, we convert input tokens w_{i-1}, \dots, w_{i-k} into embedding vectors $\{v_{i-t}\}_{t=1}^k$ and concatenate them into a dk vector $v_{i-k:i-1}$. We then pass this vector into a multilayer perceptron network with a softmax output into $|\mathcal{V}|$ classes. We implement this efficiently across a batch by using a convolution operation, since convolution acts like a moving window of size k . This way we can generate $T - k + 1$ predictions for a sentence of length T . We depict the convolution trick in [Figure 1](#).

3.3 LSTM

A concern with our trigram model is that it completely ignores words more than two positions before the word we wish to predict. To the extent we believe these words are predictive (personal experience with language suggest that they should be!), the trigram model has an inherent limitation in its ability to model that dependence.




Figure 1: Diagram for *Neural Network Language Model*

One way to combat this is with an LSTM, the architecture of which is depicted in Figure 2. At a high level, the LSTM functions by keeping track of three vectors: v_t , C_t , and h_t . The first of these vectors is simply a dense embedding for the word w_t . The h_t and C_t vectors are state representations of the model, which are dependent on previous words and give the model a “memory.” The LSTM can thus theoretically condition on all previous text, and in practice exhibits long-term memory through its architecture on C_t which encourages only intentional changes over time.

The LSTM is formally characterized by the following equations which determine the evolution of these vectors,

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) \quad (2)$$

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \quad (3)$$

$$\tilde{C}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c) \quad (4)$$

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \quad (5)$$

$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o) \quad (6)$$

$$h_t = o_t \odot \tanh(C_t) \quad (7)$$

Roughly speaking, their intuitions are as follows: \tilde{C}_t represents the new information that might be relevant for encoding into long-term memory. f_t captures the information that needs to be deleted from long-term memory, and i_t captures the places where information needs to be added. Then, these are combined to determine the new C_t . The hidden state roughly captures a filtered version (captured by the multiplication with o_t) of the cell-state.




Figure 2: Depiction of LSTM inner-workings (Olah, 2015)

3.4 Multi-head Attention

Following Vaswani et al. (2017), we implement a variant of the multi-head attention decoder network. Instead of passing the last hidden state from an LSTM h_t to predict h_{t+1} , we use a concatenation of h_t and a *context vector* $c_t = a_1 h_1 + \dots + a_{t-1} h_{t-1}$ for *attention values* a_1, \dots, a_{t-1} residing in the unit simplex. Following

Vaswani et al. (2017), we use the *scaled attention* mechanism, where

$$\mathbf{a} = \text{Softmax} \left(\left\{ \frac{h_i^T h_t}{\sqrt{\dim(h_t)}} \right\}_{i=1}^{t-1} \right).$$

In *multi-head attention*, we repeat the attention mechanism above on different linear projections of h_1, \dots, h_t , with the motivation being that we wish to capture similarity on different projections of words—one attention layer could be capturing grammar, another semantics, a third pronoun association, etc. We depict the architecture in [Figure 3](#), for $t = 3$ and predicting $t + 1$.

Certain computational tricks need to be employed for efficient utilization of the GPU. Unlike the encoding attention network, the decoder cannot condition on future information when predicting the future. As a result, each attention layer can only look at past hidden states. We parallelize the procedure and take advantage of GPU hardware by applying attention as usual, computing every inner product $h_i^T h_j$ for all i, j , and use a mask that sets entries with $h_i^T h_j$ to $-\infty$ if $j \leq i$ (which correspond to the forbidden attention links by looking ahead) before applying the softmax.

4 Experiments

5 Conclusion

References

- Bengio, Y., Ducharme, R., Vincent, P., and Jauvin, C. (2003). A neural probabilistic language model. *Journal of machine learning research*, 3(Feb):1137–1155.
- Olah, C. (2015). Understanding LSTM networks.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems*, pages 5998–6008.

A Model implementation

Listing 1: Trigram model implementation

```
1 import torch
2 from torch import Tensor
3 from torch import sparse as sp
4 from tqdm import tqdm_notebook as tqdm
5 import pandas as pd
6 import numpy as np
7 from torch import nn
8
9 def sparse_select(dims, indices, t):
10     """
11     Select sparse tensor t on dimensions dims and indices chosen by indices.
12     Equivalent to t[i_0, i_1, ...] where i_d = : if d is not in dims and
13     i_(dims[j]) = indices[j] for all j
14     """
15     if type(dims) is not list:
16         dims = [dims]
17     if type(indices) is not list:
18         indices = [indices]
19
20     t_indices = t._indices()
21     t_values = t._values()
22     selector = torch.ones(t_indices.shape[-1]).byte()
23     for dim, index in zip(dims, indices):
24         selector = selector & (t._indices()[dim, :] == index)
25     remaining_dimensions = list(filter(lambda x: x not in dims,
26                                     range(t_indices.shape[0])))
27
28     indices_selected = t_indices[:, selector][remaining_dimensions, :]
29     values_selected = t_values[selector]
30     new_shape = torch.Size(t.shape[d] for d in remaining_dimensions)
31
32     out = sp.FloatTensor(indices_selected, values_selected, new_shape)
33     return out
34
35
36
37 class Trigram:
38     def __init__(self, TEXT):
39         self.log_weights = torch.zeros(3, requires_grad=True)
40         self.TEXT = TEXT
41         self.V = len(TEXT.vocab)
42
43     def get_probabilities(self, train_iter):
44         TEXT = self.TEXT
45         V = self.V
46         unigram = sp.FloatTensor(V)
47         bigram = sp.FloatTensor(V,V)
48         trigram = sp.FloatTensor(V,V,V)
49
50         for batch in tqdm(train_iter):
51
52             i = batch.text.values.flatten().unsqueeze(0)
53             unigram_counts = sp.FloatTensor(
54                 i, torch.ones(i.shape[1]), torch.Size([V])
```

```

55         )
56
57         unigram += unigram_counts
58
59         ii = torch.stack([batch.text.values[:-1,:], batch.text.values[1:,
60             :]]).view(2, -1)
61         bigram_counts = sp.FloatTensor(
62             ii, torch.ones(ii.shape[-1]), torch.Size([V, V])
63         )
64         bigram += bigram_counts
65
66         iii = torch.stack([batch.text.values[:-2,:], batch.text.values[1:-1, :],
67             batch.text.values[2:, :]]).view(3, -1)
68         trigram_counts = sp.FloatTensor(
69             iii, torch.ones(iii.shape[-1]), torch.Size([V, V, V])
70         )
71         trigram += trigram_counts
72
73         unigram = unigram.coalesce()
74         unigram = unigram / sp.sum(unigram)
75
76         bigram = bigram.coalesce()
77         trigram = trigram.coalesce()
78
79         bigram_df = pd.DataFrame(np.hstack([bigram.indices().numpy().T,
80             bigram.values().numpy()[:, np.newaxis]]),
81             dtype=int, columns=['word1', 'word2', 'counts'])
82
83         trigram_df = pd.DataFrame(np.hstack([trigram.indices().numpy().T,
84             trigram.values().numpy()[:, np.newaxis]]),
85             dtype=int, columns=['word1', 'word2', 'word3', 'counts'])
86
87         bigram_df['prob'] = ((bigram_df['counts'] / bigram_df.groupby(['word1'])
88             .transform('sum')['counts']))
89
90         bigram_ind = torch.from_numpy(bigram_df[['word1', 'word2']].values.T)
91         bigram_val = torch.from_numpy(bigram_df['prob'].values)
92         bigram = torch.sparse.FloatTensor(bigram_ind, bigram_val, bigram.shape)
93
94         trigram_df['prob'] = (trigram_df['counts'] / trigram_df.groupby(['word1',
95             'word2']))
96             .transform('sum')['counts'])
97
98         trigram_ind = torch.from_numpy(trigram_df[['word1', 'word2',
99             'word3']].values.T)
100         trigram_val = torch.from_numpy(trigram_df['prob'].values)
101         trigram = torch.sparse.FloatTensor(trigram_ind, trigram_val, trigram.shape)
102
103         self.unigram = unigram.float()
104         self.bigram = bigram.float()
105         self.trigram = trigram.float()
106
107     def predict(self, past_two_words):
108         weights = torch.softmax(self.log_weights, dim=0)
109         output_batch = torch.zeros(len(past_two_words), self.V)
110         for i, pair in enumerate(past_two_words):
111             bi = sparse_select(0, pair[-1], self.bigram)
112             tri = sparse_select([0,1], pair.tolist(), self.trigram)

```



```

110         uni = self.unigram
111         output_batch[i, :] = (
112             tri.to_dense() * weights[0]
113             + bi.to_dense() * weights[1]
114             + uni.to_dense() * weights[2])
115         return output_batch
116
117     def __call__(self, batch_text):
118         packaged = torch.stack([batch_text.values[: -1, :], batch_text.values[1:,
119                                     :]]).view(2, -1).t()
120         return self.predict(packaged)
121
122 cross_entropy_loss = nn.CrossEntropyLoss()
123 def trigram_loss_fn(model, batch):
124     pred = model(batch.text.values)
125     labels = batch.target[1:, :].flatten()
126     loss = cross_entropy_loss(pred, labels)
127     return loss

```

Listing 2: NNLM

```

1  import torch
2  from torch import nn
3  import namedtensor
4  from namedtensor.nn import nn as namednn
5
6
7  # embedding_dim = 20
8  # hidden = 100
9  # seqlen = 32
10 # pad_i = TEXT.vocab.stoi['<pad>']
11
12 # class Pad(nn.Module):
13 # def __init__(self, seqlen, pad_i):
14 # super().__init__()
15 # self.seqlen = seqlen
16 # self.pad_i = pad_i
17 # def forward(self, x):
18 # init = torch.ones(self.seqlen, x.shape[1]) * pad_i
19 # init[-x.shape[0]:, :] = x
20 # return init.long()
21
22
23 # class Flatten(nn.Module):
24 # def __init__(self):
25 # super().__init__()
26 # def forward(self, x):
27 # return x.permute(1,0,2).flatten(start_dim=1,end_dim=2)
28
29 # def get_nn_lang_model(embedding_dim, hidden, TEXT, n_hidden_layers=1, seqlen=32):
30 # pad_i = TEXT.vocab.stoi['<pad>']
31 # if n_hidden_layers == 1:
32 # return nn.Sequential(
33 #     Pad(seqlen, pad_i),
34 #     nn.Embedding(len(TEXT.vocab), embedding_dim),
35 #     Flatten(),
36 #     nn.Linear(seqlen * embedding_dim, hidden),
37 #     nn.Tanh(),
38 #     nn.Linear(hidden, len(TEXT.vocab)),

```

```

39 # nn.Softmax(dim=0)
40 # )
41 # else:
42 # layers = []
43 # for k in range(n_hidden_layers-1):
44 # layers.append(nn.Linear(hidden, hidden))
45 # layers.append(nn.Tanh())
46 # return nn.Sequential(
47 # Pad(seqlen, pad_i),
48 # nn.Embedding(len(TEXT.vocab), embedding_dim),
49 # Flatten(),
50 # nn.Linear(seqlen * embedding_dim, hidden),
51 # nn.Tanh(),
52 # *layers,
53 # nn.Linear(hidden, len(TEXT.vocab)),
54 # nn.Softmax(dim=0)
55 # )
56
57
58 class NNLangModel(namednn.Module):
59     def __init__(self, TEXT, embedding_dim, kernel_size, hidden, dropout=.5):
60         super().__init__()
61         V = len(TEXT.vocab)
62         pad_idx = TEXT.vocab.stoi['<pad>']
63
64         self.embed = namednn.Embedding(num_embeddings=V,
65                                         embedding_dim=embedding_dim,
66                                         padding_idx=pad_idx)
67         self.conv = namednn.Conv1d(embedding_dim, embedding_dim,
68                                     kernel_size=kernel_size).spec('embedding', 'seqlen')
69
70         self.w1 = namednn.Linear(embedding_dim, hidden).spec('embedding', 'hidden')
71         self.w2 = namednn.Linear(hidden, hidden).spec('hidden', 'hidden2')
72         self.w3 = namednn.Linear(hidden, V).spec('hidden2', 'classes')
73         self.dropout = namednn.Dropout(dropout)
74
75     def forward(self, batch_text):
76         embedded = self.embed(batch_text)
77         conved = self.conv(embedded)
78         h1 = self.w1(conved).tanh()
79         h2 = self.w2(self.dropout(h1)).tanh()
80         out = self.w3(self.dropout(h2))
81         return out
82
83 nn_lang_loss = namednn.CrossEntropyLoss().spec('seqlen')
84 def nn_lang_loss_fn(model, batch):
85     output = model(batch.text)
86     size = output.shape['seqlen']
87     target_size = batch.target.size('seqlen')
88     target = (batch.target[{'seqlen' : slice(target_size-size, target_size)}]
89               .transpose('batch', 'seqlen'))
90     return nn_lang_loss(output, target)

```

Listing 3: LSTM

```

1 from namedtensor.nn import nn as nnn
2
3 class LSTM(nnn.Module):
4     """

```

```

5 LSTM implementation for sentence completion.
6 """
7 def __init__(self, TEXT,
8               embedding_dim=100,
9               hidden_dim=150,
10              num_layers=1,
11              dropout=0):
12     super().__init__()
13
14     pad_idx = TEXT.vocab.stoi['<pad>']
15
16     self.embed = nn.Embedding(num_embeddings=len(TEXT.vocab),
17                              embedding_dim=embedding_dim,
18                              padding_idx=pad_idx)
19
20     self.lstm = nn.LSTM(input_size=embedding_dim,
21                        hidden_size=hidden_dim,
22                        num_layers=num_layers,
23                        dropout=dropout) \
24         .spec("embedding", "seqlen")
25
26     self.w = nn.Linear(in_features=hidden_dim,
27                       out_features=len(TEXT.vocab)) \
28         .spec("embedding", "classes")
29
30     def forward(self, batch_text):
31         embedded = self.embed(batch_text)
32         hidden_states, _ = self.lstm(embedded)
33         log_probs = self.w(hidden_states)
34
35         return log_probs
36
37
38 ce_loss = nn.CrossEntropyLoss().spec('batch')
39
40 def lstm_loss(model, batch):
41     """
42     Calculate loss of the model on a batch.
43     """
44     return ce_loss(model(batch.text), batch.target)

```

Listing 4: LSTM-attention

```

1 from namedtensor.nn import nn as nn
2 from namedtensor import ntorch
3 import torch
4 from namedtensor import NamedTensor
5 from numpy import inf
6
7
8 class MaskedAttention(nn.Module):
9     def __init__(self, cuda=True):
10         super().__init__()
11         self.cuda_enabled = cuda
12
13     def forward(self, hidden):
14         dotted = (hidden * hidden.rename("seqlen", "seqlen2")).sum("embedding")
15         mask = torch.arange(hidden.size('seqlen'))
16         mask = (NamedTensor(mask, names='seqlen') < NamedTensor(mask,

```

```

17         names='seqlen2')).float()
18     mask[mask.byte()] = -inf
19     if self.cuda_enabled:
20         attn = ((dotted + mask.cuda()) / (hidden.size("embedding") **
21             .5)).softmax('seqlen2')
22     else:
23         attn = ((dotted + mask) / (hidden.size("embedding") **
24             .5)).softmax('seqlen2')
25     return (attn * hidden.rename('seqlen', 'seqlen2')).sum('seqlen2')
26
27 class LSTM_att(nnn.Module):
28     """
29     LSTM implementation for sentence completion.
30     """
31     def __init__(self, TEXT,
32         embedding_dim=100,
33         hidden_dim=150,
34         num_layers=1,
35         dropout=0,
36         nn_dropout=.5,
37         **kwargs):
38         super().__init__()
39
40         pad_idx = TEXT.vocab.stoi['<pad>']
41
42         self.embed = nnn.Embedding(num_embeddings=len(TEXT.vocab),
43             embedding_dim=embedding_dim,
44             padding_idx=pad_idx)
45
46         self.lstm = nnn.LSTM(input_size=embedding_dim,
47             hidden_size=hidden_dim,
48             num_layers=num_layers,
49             dropout=dropout) \
50             .spec("embedding", "seqlen")
51
52         self.w1 = (nnn.Linear(in_features=hidden_dim, out_features=hidden_dim)
53             .spec("embedding", "embedding"))
54         self.w2 = (nnn.Linear(in_features=hidden_dim, out_features=hidden_dim)
55             .spec("embedding", "embedding"))
56         self.w3 = (nnn.Linear(in_features=hidden_dim, out_features=hidden_dim)
57             .spec("embedding", "embedding"))
58
59         self.lins = [self.w1, self.w2, self.w3]
60         self.attn = MaskedAttention(**kwargs)
61
62         h_len = len(self.lins) + 2
63         self.w = nnn.Linear(in_features=hidden_dim * h_len,
64             out_features=len(TEXT.vocab)) \
65             .spec("embedding", "classes")
66         self.dropout = nnn.Dropout(nn_dropout)
67
68     def forward(self, batch_text):
69         embedded = self.embed(batch_text)
70         H, _ = self.lstm(embedded)
71         joint = ntorch.cat([H, self.attn(H)] + [self.attn(l(H)) for l in self.lins],
72             "embedding")

```

```

72     log_probs = self.w(self.dropout(joint))
73     return log_probs
74
75
76 ce_loss = nn.CrossEntropyLoss().spec('classes')
77
78 def lstm_loss(model, batch):
79     """
80     Calculate loss of the model on a batch.
81     """
82     return ce_loss(model(batch.text), batch.target)

```

Listing 5: Ensemble

```

1  from namedtensor import ntorch
2
3  class Ensemble:
4      def __init__(self, *models):
5          self.models = models
6
7      def __call__(self, batch_text):
8          return ntorch.stack(
9              [model(batch_text) for model in self.models], "model").mean("model")

```




Figure 3: Diagram for *Multi-head Attention*. In this diagram, we predict the fourth word in the sentence by conditioning on the first three. We compute attention values of h_3 with h_2 and h_1 and concatenate h_3 with a context vector $a_1h_1 + a_2h_2$, where a_i are the attention values. We pass the resulting vector through a linear transformation and softmax output. In multi-head attention, we repeat the process for different projections of h_1, h_2, h_3 and concatenate the results.